

ECE/CS 472/572:
Special Topics Part III:
GPU Architecture Introduction

Prof. Lizhong Chen

Spring 2019


GPU: System Organization

- Nvidia DGX-1 High-performance computing (HPC) sever




GPU: System Organization

- Nvidia DGX-1 with panel open




GPU: System Organization

- Inside Nvidia DGX-1



GPU: System Organization

- Tesla P100 in Nvidia DGX-1



NVIDIA Tesla GP100



- 6 GPCs
- 30 TPCs
- 60 SMs
- 3840 Cores
- 240 Tex
- 8 MCs
 - 4096-bit


NVIDIA Telsa GP100 SM



- 64 FP32 cores
 - Supports 16-bit FP
- 32 FP64 cores
- 2 32K 32-bit Registers
- Separate shared mem.
- Unified Texture/L1 \$


GPUs Everywhere

- Many of the top-ranked supercomputers are based on GPU
 - The World #1 supercomputer is powered by NVIDIA GPU



But Why are GPUs Everywhere?

- It is the performance/Watt metric
 - Mobile computers, Supercomputers, Laptops: Share power agony



GPU: Initially for Graphics

- Take a scene and project on a screen
 - How does each polygon translate into a screen pixel





Image credit: Henrik Wann Jensen


Graphics-based GPU Pipeline (before 2006)





CS 550: Introduction to Computer Graphics

GPU Pipeline Functions

- **Vector Shader** : Takes as input vectors given in 3D format (X,Y,Z) and projects to a 2D space
- **Pixel Shader**: Takes each projected point in 2D space and puts texture, color and depth for each point
 - Texture mapping: Adding an image on top of a scene
 - Depth mapping: Shows which objects will obscure other objects



Example Image: Before and After Texture Mapping



Example Image: Z-Buffer Representation of Image



- **Geometry Shaders**: Operates directly on graphics primitives, such as lines, triangles (as opposed to pixels)

CS 550: Introduction to Computer Graphics

Need for Unified Shader

Why unify?

Unified Shader



Unified Shader



© NVIDIA Corporation 2007



**Heavy Geometry
Workload Perf = 11**



**Heavy Pixel
Workload Perf = 11**

Resource Balancing Unified Shader

Dynamic Load Balancing – Company of Heroes



Less Geometry



More Geometry




Unified Shader

GPU Pipeline evolution


Before 2006

- Processor per function

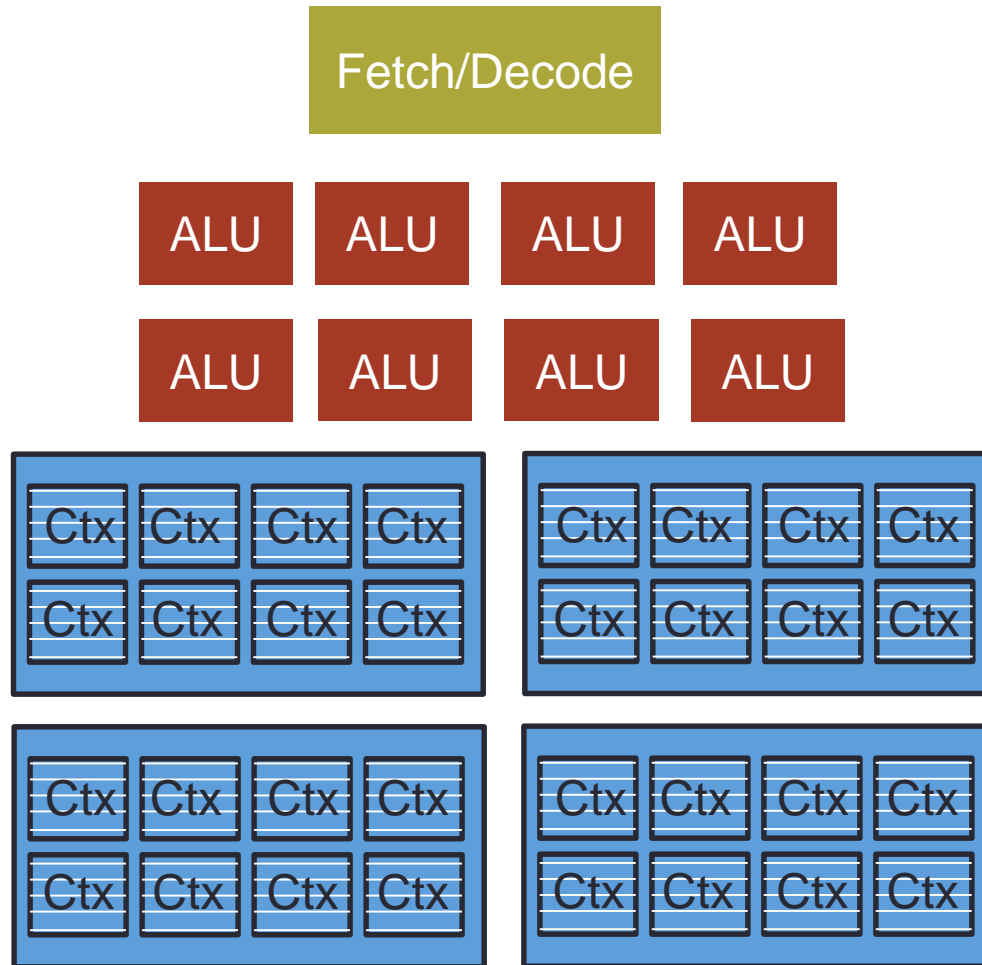


After 2006


- Unified Shader Model(GeForce 8 Series)
- Any work can be performed on any shader core → High performance computing




Unified Shader Model



Unified Shader Model

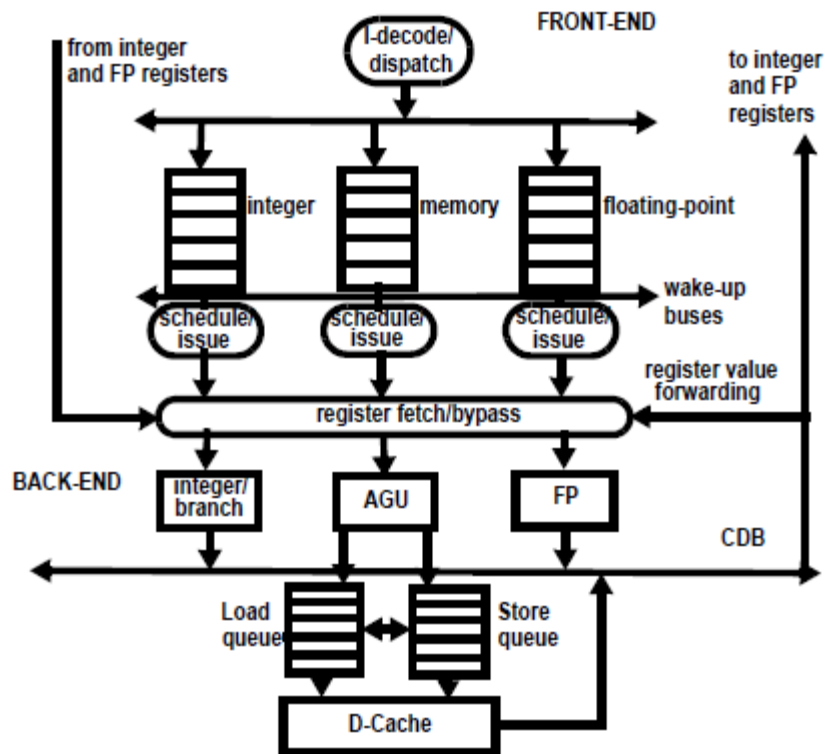


Multiple Contexts(Thread groups)
can be interleavingly scheduled
→ Can hide stall times




GPU vs CPU

Complex OoO CPU



GPU (Nvidia FERMI)



Basic idea of high performance GPU

- Many simple cores
 - No fancy Branch prediction
 - No Complex O-o-O control logic
 - No memory prefetcher
 - No cache coherence
 - ...
- Unified shader cores
 - Sharing instruction stream across groups of fragments
- Stall latency hiding
 - When a group stalls, work on another group